

+ Code + Text

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[100] df=pd.read_csv("/content/sample_data/winequalityN.csv")
df
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	white	7.0	0.270	0.36	20.7	0.045	45.0	170.0	1.00100	3.00	0.45	8.8	6
1	white	6.3	0.300	0.34	1.6	0.049	14.0	132.0	0.99400	3.30	0.49	9.5	6
2	white	8.1	0.280	0.40	6.9	0.050	30.0	97.0	0.99510	3.26	0.44	10.1	6
3	white	7.2	0.230	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
4	white	7.2	0.230	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
...
6492	red	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5

1s completed at 9:12 AM

`df.describe()` 0s

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
count	6487.000000	6489.000000	6494.000000	6495.000000	6495.000000	6497.000000	6497.000000	6497.000000	6488.000000	6493.000000	6497.000000	6497.000000
mean	7.216579	0.339691	0.318722	5.444326	0.056042	30.525319	115.744574	0.994697	3.218395	0.531215	10.491801	5.87
std	1.296750	0.164649	0.145265	4.758125	0.035036	17.749400	56.521855	0.002999	0.160748	0.148814	1.192712	0.87
min	3.800000	0.080000	0.000000	0.600000	0.009000	1.000000	6.000000	0.987110	2.720000	0.220000	8.000000	3.00
25%	6.400000	0.230000	0.250000	1.800000	0.038000	17.000000	77.000000	0.992340	3.110000	0.430000	9.500000	5.00
50%	7.000000	0.290000	0.310000	3.000000	0.047000	29.000000	118.000000	0.994890	3.210000	0.510000	10.300000	6.00
75%	7.700000	0.400000	0.390000	8.100000	0.065000	41.000000	156.000000	0.996990	3.320000	0.600000	11.300000	6.00
max	15.900000	1.580000	1.660000	65.800000	0.611000	289.000000	440.000000	1.038980	4.010000	2.000000	14.900000	9.00





+ Code + Text

 RAM
Disk

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   type                  6497 non-null   object
1   fixed acidity          6487 non-null   float64
2   volatile acidity       6489 non-null   float64
3   citric acid            6494 non-null   float64
4   residual sugar         6495 non-null   float64
5   chlorides              6495 non-null   float64
6   free sulfur dioxide    6497 non-null   float64
7   total sulfur dioxide   6497 non-null   float64
8   density               6497 non-null   float64
9   pH                    6488 non-null   float64
10  sulphates              6493 non-null   float64
11  alcohol                6497 non-null   float64
12  quality                6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

1s completed at 9:12 AM





+ Code + Text

```
[104] 12 quality          6497 non-null    int64  
0s dtypes: float64(11), int64(1), object(1)  
memory usage: 660.0+ KB
```

```
{x} df.isna().sum()  
0s
```

```
type          0  
fixed acidity  10  
volatile acidity    8  
citric acid        3  
residual sugar     2  
chlorides          2  
free sulfur dioxide  0  
total sulfur dioxide  0  
density           0  
pH                9  
sulphates         4  
alcohol           0  
quality           0  
dtype: int64
```


```
[106] # create dist plot  
6s import warnings  
%matplotlib inline  
warnings.filterwarnings('ignore')
```

1s completed at 9:12 AM

```
# create dist plot
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')

fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

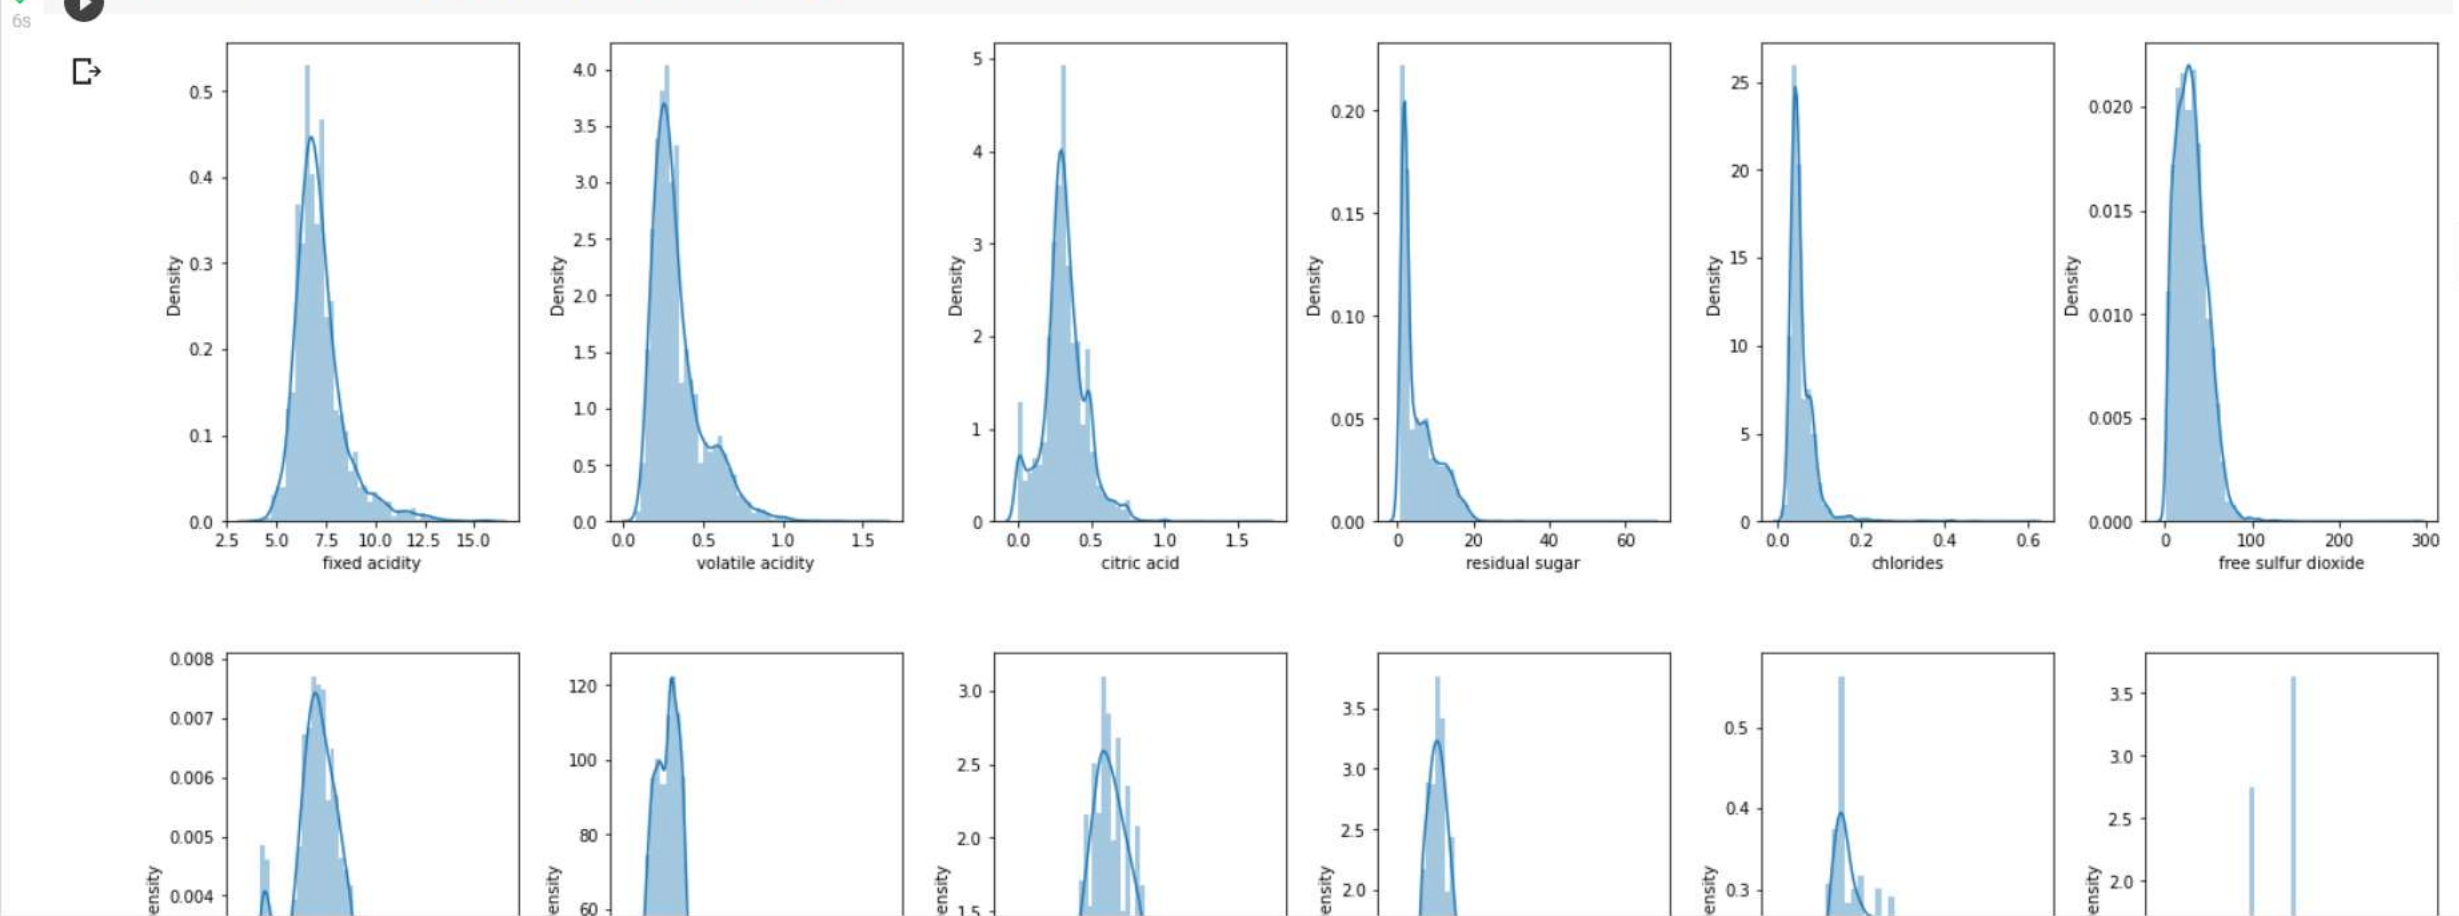
for col, value in df.items():
    if col != 'type':
        sns.distplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



✓ 1s completed at 9:12 AM ● ✕

+ Code + Text


```
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



✓ 1s completed at 9:12 AM



+ Code + Text

 RAM
Disk 0s `sns.countplot('type',data=df)` `<matplotlib.axes._subplots.AxesSubplot at 0x7f890ee4cac0>` 2s `[108] sns.jointplot(x='free sulfur dioxide',y='total sulfur dioxide',data=df,kind='reg')``<seaborn.axisgrid.JointGrid at 0x7f890d0b30d0>`

1s completed at 9:12 AM




+ Code + Text

RAM Disk


`sns.jointplot(x='free sulfur dioxide',y='total sulfur dioxide',data=df,kind='reg')`

2s

`<seaborn.axisgrid.JointGrid at 0x7f890d0b30d0>`


1s completed at 9:12 AM

Code Text



✓ 0s `sns.distplot(df['alcohol'])`

`<matplotlib.axes._subplots.AxesSubplot at 0x7f894e0ea910>`




✓ [110] 0s `sns.countplot(df['quality'])`

`<matplotlib.axes._subplots.AxesSubplot at 0x7f8909ebc940>`




✓ 1s completed at 9:12 AM

Code Text



✓ 0s `sns.countplot(df['quality'])`

`<matplotlib.axes._subplots.AxesSubplot at 0x7f8909ebc940>`




[111] `#coorelation matrix`
`corr=df.corr()`
`plt.figure(figsize=(20,10))`

✓ 1s completed at 9:12 AM

+ Code + Text

```
#coorelation matrix
corr=df.corr()
plt.figure(figsize=(20,10))
sns.heatmap(corr,annot=True,cmap='coolwarm')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f890d0f2eb0>




1s completed at 9:12 AM

+ Code + Text

✓ RAM
Disk

👤 ⚙️ ⌵

[111]

✓
0s

df.isna().sum()

```
type                0
fixed acidity       10
volatile acidity     8
citric acid         3
residual sugar      2
chlorides           2
free sulfur dioxide  0
total sulfur dioxide 0
density             0
pH                 9
sulphates           4
alcohol             0
quality             0
dtype: int64
```

✓ 1s completed at 9:12 AM

● x

+ Code + Text

✓ RAM
Disk

👤 ⚙️ ▾

```
[113] df['fixed acidity']=df['fixed acidity'].fillna(df['fixed acidity'].mean())
      df['volatile acidity']=df['volatile acidity'].fillna(df['volatile acidity'].mean())
      df['citric acid']=df['citric acid'].fillna(df['citric acid'].mean())
      df['residual sugar']=df['residual sugar'].fillna(df['residual sugar'].mean())
      df['chlorides']=df['chlorides'].fillna(df['chlorides'].mean())
      df['pH']=df['pH'].fillna(df['pH'].mean())
      df['sulphates']=df['sulphates'].fillna(df['sulphates'].mean())
```

```
df.isna().sum()
```

```
type           0
fixed acidity   0
volatile acidity 0
citric acid     0
residual sugar  0
chlorides       0
free sulfur dioxide 0
total sulfur dioxide 0
density         0
pH              0
sulphates       0
alcohol         0
quality         0
dtype: int64
```

```
[115] df= df.drop(['type'],axis =1)
```

✓ 1s completed at 9:12 AM

● x

```
df= df.drop(['type'],axis =1)
df
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.0	0.270	0.36	20.7	0.045	45.0	170.0	1.00100	3.00	0.450000	8.8	6
1	6.3	0.300	0.34	1.6	0.049	14.0	132.0	0.99400	3.30	0.490000	9.5	6
2	8.1	0.280	0.40	6.9	0.050	30.0	97.0	0.99510	3.26	0.440000	10.1	6
3	7.2	0.230	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.400000	9.9	6
4	7.2	0.230	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.400000	9.9	6
...
6492	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.580000	10.5	5
6493	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.531215	11.2	6
6494	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.750000	11.0	6
6495	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.710000	10.2	5
6496	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.660000	11.0	6

6497 rows x 12 columns



+ Code + Text

✓ RAM
Disk

```
[116] x=df.iloc[:, :-1].values  
      y=df.iloc[:, -1].values  
      y
```

```
array([6, 6, 6, ..., 6, 5, 6])
```

```
[117] from sklearn.model_selection import train_test_split  
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=42)  
      y_test
```

```
array([7, 7, 6, ..., 6, 6, 6])
```

```
from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()  
scaler.fit(x_train)  
x_train=scaler.transform(x_train)  
x_test=scaler.transform(x_test)  
x_test
```

```
array([[ -0.17304369, -0.54359665,  0.90004528, ..., -0.36283779,  
        -0.35375297,  1.17959322],  
       [ 0.28607349, -1.20786028,  2.89927908, ..., -0.92242828,  
        -0.88352406,  0.256953  ],  
       [-0.78519994, -1.14747268, -0.34085846, ..., -0.61154467,  
        -1.01596683, -1.16894553],  
       ...,  
       ...])
```

✓ 1s completed at 9:12 AM

● X

76°F
Sunny

Search

ENG
IN09:20
08-02-2023

2

+ Code + Text

✓ RAM
Disk

👤 ⚙️ ▼

```
[119] from sklearn.tree import DecisionTreeClassifier
      model=DecisionTreeClassifier(criterion='entropy')
      model.fit(x_train,y_train)
      y_pred=model.predict(x_test)
      y_pred
```

```
array([7, 7, 6, ..., 5, 6, 6])
```

```
df1=pd.DataFrame({'Actual_Value':y_test,'Predicted_Value':y_pred})
df1
```

📄 Actual_Value Predicted_Value ✎

0	7	7
1	7	7
2	6	6
3	6	6
4	5	5
...
1295	7	7
1296	5	6
1297	6	5

✓ 1s completed at 9:12 AM

● x

+ Code + Text

✓ RAM  Disk  👤 ⚙️ ▼

```
[121] from sklearn.metrics import classification_report, accuracy_score, confusion_matrix, ConfusionMatrixDisplay
      result=confusion_matrix(y_test,y_pred)
      score=accuracy_score(y_test,y_pred)
      score
```

0.6138461538461538

```
▶ from sklearn.neighbors import KNeighborsClassifier
  classifier=KNeighborsClassifier(n_neighbors=7)
  classifier.fit(x_train,y_train)
  y_pred=classifier.predict(x_test)
  print(y_pred)
```

[7 6 6 ... 5 6 6]

```
[123] from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
      result=confusion_matrix(y_test,y_pred)
      score=accuracy_score(y_test,y_pred)
      score
```

0.5553846153846154

```
[125] from sklearn.svm import SVC
      classifier=SVC()
      classifier.fit(x_train,y_train)
      v_pred=classifier.predict(x_test)
```

✓ 1s completed at 9:12 AM

● X

+ Code + Text

✓ RAM
Disk

```
[125] from sklearn.svm import SVC
      classifier=SVC()
      classifier.fit(x_train,y_train)
      y_pred=classifier.predict(x_test)
      print(y_pred)
```

```
[6 6 6 ... 5 6 6]
```

```
✓ [126] from sklearn.metrics import classification_report,accuracy_score,confusion_matrix,ConfusionMatrixDisplay
0s      result=confusion_matrix(y_test,y_pred)
      score=accuracy_score(y_test,y_pred)
      score
```

```
0.5815384615384616
```

```
✓ 2s ▶ from sklearn.ensemble import RandomForestClassifier
      clf = RandomForestClassifier(n_estimators = 100)
      clf.fit(x_train, y_train)
      y_pred = clf.predict(x_test)

      from sklearn import metrics
      print()

      print("Accuracy_Score: ", metrics.accuracy_score(y_test, y_pred))
```

```
Accuracy_Score: 0.6907692307692308
```

✓ 1s completed at 9:12 AM

● X